Applied Business Research

Eric Weisbrod, with collaboration from ACCT 995 Students

2025-01-22

Table of contents

Preface

This is a Quarto book.

To learn more about Quarto books visit https://quarto.org/docs/books.

```
1 + 1
```

```
[1] 2
```

1 Introduction

This is a book created from markdown and executable code.

See Knuth (1984) for additional discussion of literate programming.

```
1 + 1
```

```
[1] 2
```

2 Obtaining and Merging Data

This is my first time working on a Quarto book. So, this first post will be very rough for now. I will try to provide a few different examples of ways to obtain and merge data in R, and a few tips of things to keep in mind.

We already know how to obtain data from WRDS. Let's use this to obtain some returns for the S&P 500. We could use the formal index data, but let's take a shortcut and just use the popular SPY ETF that tracks the S&P 500. To do this, we need to find the CRSP identifier (PERMNO) for the ticker "SPY." We can look in the WRDS stocknames file for this, and then use the SPY PERMNO to pull data from the CRSP monthly stock file.

```r
# Load Libraries [i.e., packages]
library(dbplyr)
library(RPostgres)
library(DBI)
library(glue)
library(arrow)
library(haven)
library(tictoc) #very optional timer, mostly as a teaching example
library(tidyverse) # I like to load tidyverse last to avoid package conflicts

#I have done this in a separate chunk with the options
# results: FALSE
# message: FALSE
#because I don't need to see the messages from loading the packages.
```

```r
# Log in to WRDS ----------------------------------------------------------------

#before running this block, I used these commands to securely store my WRDS username and pass
# keyring::key_set("WRDS_user")
# keyring::key_set("WRDS_pw")

if(exists("wrds")){
  dbDisconnect(wrds)  # because otherwise WRDS might time out
}
```

```
wrds <- dbConnect(Postgres(),
                  host='wrds-pgdata.wharton.upenn.edu',
                  port=9737,
                  user=keyring::key_get("WRDS_user"),
                  password=keyring::key_get("WRDS_pw"),
                  sslmode='require',
                  dbname='wrds')



# Create WRDS Table References ----------------------------------------
crsp.msf <- tbl(wrds,in_schema("crsp","msf"))
stocknames <- tbl(wrds,in_schema("crsp","stocknames"))

#I am collecting this data locally to play with duplicates
spy_permnos <- stocknames |> filter(ticker == "SPY") |> collect()
```

Notice that there are six observations in the stocknames table that all share the same ticker "SPY." I am going to use this as a toy example to play with duplicates. My goal is for this data to be unique at the level of ticker-permno links. First, I can check whether this is true.

```
#check whether there are duplicates
#this simple logic is useful in general
#group by the level I want to make unique,
#count within each group
#sort by descending count so that if there are duplicates
#they will show up at the top.
spy_permnos |>
  group_by(ticker,permno) |>
  count() |>
  arrange(-n)
```

```
# A tibble: 3 x 3
# Groups:   ticker, permno [3]
  ticker permno     n
  <chr>   <int> <int>
1 SPY     84398     3
2 SPY     33910     1
3 SPY     60716     1
```

There are multiple permnos connected to the SPY ticker and some duplicate entries for permno 84398 so I better just look at the data. Also this tells me that there are only a few rows so it doesn't hurt to just print the data.

```
#| #note that we can use the kable commmand to embed a simple table in the quarto document
knitr::kable(spy_permnos)
```

permno	namedt	nameenddt	shrcd	exchcd	siccd	ncusip	ticker	comnam	shrcls	permno	hexcd	cusip	st_date	end_date	namedum
33910	1962-07-02	1966-05-24	10	2	2893	NA	SPY	SPEEDRY CHEMICAL PRODS INC	A	2751	3	559142	1962-07-02	1979-01-22	2
60716	1978-10-03	1987-07-01	10	1	3811	84756510	SPY	SPECTRA PHYSICS INC	NA	4215	1	847565	1972-12-14	1987-07-01	2
84398	1993-01-29	2009-02-23	73	2	6726	78462F10	SPY	SPDR TRUST	NA	46699	4	78462F	1993-01-29	2024-12-31	2
84398	2009-02-24	2010-01-26	73	4	6726	78462F10	SPY	SPDR TRUST	NA	46699	4	78462F	1993-01-29	2024-12-31	2
84398	2010-01-27	2024-12-31	73	4	6726	78462F10	SPY	SPDR S & P 500 E T F TRUST	NA	46699	4	78462F	1993-01-29	2024-12-31	2

Looking at the data, the company name for permno 84398 matches the SPDR S&P 500 ETF I am looking for. It looks like the duplicate entries might have to do with a change in the listing exchange for the ETF (exchcd) and then a slight name change in 2010 to make the name of the trust more descriptive. Let's keep using this toy example to demonstrate some other functions for dealing with duplicates:

```
#if I want to just collapse the duplicates, I can use "distinct" across the groups that I ca

spy_permnos |>
  select(ticker,permno) |>
  distinct()
```

```
# A tibble: 3 x 2
  ticker permno
  <chr>   <int>
1 SPY     33910
2 SPY     60716
3 SPY     84398
```

Now there are only three observations,which is what I asked for, but sometimes it might matter which of the duplicate observations I keep. For example, perhaps what I should do is keep the most recent observation from the spy_permno dataset, in terms of nameenddt.

```
#select the max data within each group as more advanced way to keep one obs per
#group
spy_permnos |>
  group_by(ticker,permno) |>
  filter(nameenddt==max(nameenddt))
```

```
# A tibble: 3 x 16
# Groups:   ticker, permno [3]
  permno namedt     nameenddt  shrcd exchcd siccd ncusip   ticker comnam  shrcls
   <int> <date>     <date>     <int>  <int> <int> <chr>    <chr>  <chr>   <chr>
1  33910 1962-07-02 1966-05-24    10      2  2893 <NA>     SPY    SPEEDR~ A
2  60716 1978-10-03 1987-07-01    10      1  3811 84756710 SPY    SPECTR~ <NA>
3  84398 2010-01-27 2024-12-31    73      4  6726 78462F10 SPY    SPDR S~ <NA>
# i 6 more variables: permco <int>, hexcd <int>, cusip <chr>, st_date <date>,
#   end_date <date>, namedum <dbl>
```

```
#ultimately we can assign the permno of the current observation, which we already know from r

spy_permno <- spy_permnos |>
  group_by(ticker,permno) |>
  filter(nameenddt==max(nameenddt)) |>
  ungroup() |>
  filter(nameenddt==max(nameenddt)) |>
  select(permno) |>
  as.numeric()

spy_permno
```

```
[1] 84398
```

Now we can use the SPY permno to pull monthly returns for SPY:

```
# Pull CRSP MSI Data ------------------------------------------------------

#Data seems to begin in feb 1993, lets start in 1995 as a nice round number
#notice that this implicitly feeds the permno I calculated locally back up to WRDS in my crsp
mkt_index <- crsp.msf |>
```

```
    filter(date >= "1995-01-01",
           permno == spy_permno) |>
    select(date,ret,prc) |>
    collect() |>
    mutate(month = month(date),
           year = year(date))
```

Then I can plot them, note that if you look at the source code for this page, I do this in a chunk with echo=false so that I only see the output and not the code. This would be useful for creating an actual paper rather than coding examples:

This plot would look nice with recessions shaded. We can get recession dates from FRED. FRED data can be accessed from an API, there is a custom package to work with FRED data in R called fredr. First you need to obtain a FRED API key by signing up here: https: //fred.stlouisfed.org/docs/api/api_key.html

```
#load the fredr package
library(fredr)

#set my API key which is saved in keyring
fredr_set_key(keyring::key_get("fred_api_key"))

#collect the data from the series USRECD
```

```
# https://fred.stlouisfed.org/series/USRECD

fred_data<-fredr(series_id = "USRECD",
                 observation_start = as.Date("1995-01-01"),
                 observation_end = as.Date("2024-12-31"),
                 frequency = "m") |>
  #I am going to add month and year variables because I think this is
  #easier for linking
  mutate(month = month(date),
         year = year(date))

# show the first few rows which has a value of 0 or 1 where 1 is recession
fred_data |> head() |> knitr::kable()
```

date	series_id	value	realtime_start	realtime_end	month	year
1995-01-01	USRECD	0	2025-03-07	2025-03-07	1	1995
1995-02-01	USRECD	0	2025-03-07	2025-03-07	2	1995
1995-03-01	USRECD	0	2025-03-07	2025-03-07	3	1995
1995-04-01	USRECD	0	2025-03-07	2025-03-07	4	1995
1995-05-01	USRECD	0	2025-03-07	2025-03-07	5	1995
1995-06-01	USRECD	0	2025-03-07	2025-03-07	6	1995

Now we need to merge the SPY data with the recession data.

```
merged_data <- mkt_index |>
  #I am going to select only the columns I need from   #the FRED data
  inner_join(fred_data |>
               select(month,year,recession=value),
             by=join_by(month,year))

# check to make sure it is still unique by month
merged_data |>
  group_by(month,year) |>
  count() |>
  arrange(-n)
```

```
# A tibble: 360 x 3
# Groups:   month, year [360]
   month  year     n
   <dbl> <dbl> <int>
```

```
 1      1  1995      1
 2      1  1996      1
 3      1  1997      1
 4      1  1998      1
 5      1  1999      1
 6      1  2000      1
 7      1  2001      1
 8      1  2002      1
 9      1  2003      1
10      1  2004      1
# i 350 more rows
```

Now we can make the plot with shades for recession months

```
#turns out the merged data was not the preferred way to do this kind of plot




#here is some code I found online to reshape the recession data and add it to the plot

#rename/assign fred data to recession because
#that was the name in the example I found
recession<-fred_data

#load a package they used
library(ecm)

#reshape the recession data for the way
#geom_rect likes the data shaped
recession$diff<-recession$value-lagpad(recession$value,k=1)
  recession<-recession[!is.na(recession$diff),]
  recession.start<-recession[recession$diff==1,]$date
  recession.end<-recession[recession$diff==(-1),]$date

  if(length(recession.start)>length(recession.end))
  {recession.end<-c(recession.end,Sys.Date())}
  if(length(recession.end)>length(recession.start))
  {recession.start<-c(min(recession$date),recession.start)}

  recs<-as.data.frame(cbind(recession.start,recession.end))
  recs$recession.start<-as.Date(as.numeric(recs$recession.start),origin=as.Date("1970-01-01")
```

```
    recs$recession.end<-as.Date(recs$recession.end,origin=as.Date("1970-01-01"))

#look at the reshaped data
recs
```

```
  recession.start recession.end
1      2001-04-01    2001-12-01
2      2008-01-01    2009-07-01
3      2020-03-01    2020-05-01
```

```
#plot the new plot with recession bars
merged_data |>
  ggplot(aes(x=date,y=abs(prc))) +
  geom_line() +
  scale_x_date(name = "Date",
               date_breaks= "5 years",
               date_labels = "%Y") +
  scale_y_continuous(name = "SPY Closing Price") +
  geom_rect(data=recs, inherit.aes=F,
                         aes(xmin=recession.start, xmax=recession.end, ymin=-Inf, ymax=+Inf)
              fill="darkgrey", alpha=0.5)+
  theme_bw()
```

3 Regression Tables

Test of embedding a regression in Quarto.

Table 3.1

	Base	No FE	Year FE	Two-Way FE	With Controls
ROA_t	0.839***	0.756***	0.769***	0.639***	0.624***
	(62.732)	(48.155)	(48.621)	(38.634)	(35.596)
$LOSS$		-0.030***	-0.028***	-0.015***	-0.017***
		(-7.949)	(-7.755)	(-7.556)	(-8.111)
$ROA_t \times LOSS$		0.032	0.012	-0.285***	-0.294***
		(1.470)	(0.535)	(-13.307)	(-12.620)
Year FE			X	X	X
Firm FE				X	X
Controls					X
N	163,298	163,298	163,298	161,635	161,635
R^2	0.594	0.597	0.603	0.707	0.707
R^2 Within			0.580	0.184	0.186

4 Summary

In summary, this book has no content whatsoever.

```
1 + 1
```

```
[1] 2
```

References

Knuth, Donald E. 1984. "Literate Programming." *Comput. J.* 27 (2): 97–111. https://doi.org/10.1093/comjnl/27.2.97.